

# Graph Databases

Lecture plan

1. Many to many relationships
2. Graph databases
3. Neo4j

## 1 Many to many relationships

So far most of the data we have dealt with has had **one to one** or **one to many** relationships

Some data does not fit into these models and instead have **many to many** relationships

❓ Can you **think of any examples?**

## 🎯 Lets try modeling a many to many relationships in SQL

It will be a user and friends relationship here most users will have many friends on the platform which are other users!

Lets create a table `users` and a table `friends`

```
CREATE TABLE users (
    id SERIAL PRIMARY KEY,
    name VARCHAR(100) NOT NULL,
    email VARCHAR(100) NOT NULL UNIQUE
);
CREATE TABLE friends (
    user_id INTEGER REFERENCES users(id),
    friend_id INTEGER REFERENCES users(id),
    PRIMARY KEY (user_id, friend_id)
);
```

Lets populate `users` with 10000 users

```

DO $$

DECLARE
    counter INTEGER := 0;
BEGIN
    WHILE counter < 10000 LOOP
        INSERT INTO users (id, name, email)
        VALUES (
            counter,
            'User ' || counter,
            'user' || counter || '@example.com'
        )
        ON CONFLICT (id) DO NOTHING;
        counter := counter + 1;
    END LOOP;
END $$;

```

Lets populate friends with 20000 random friendships!

```

DO $$

DECLARE
    counter INTEGER := 0;
    user1 INTEGER;
    user2 INTEGER;
    max_id INTEGER;
BEGIN
    SELECT INTO max_id MAX(id) FROM users;
    WHILE counter < 20000 LOOP
        user1 := FLOOR(RANDOM() * max_id) + 1;
        user2 := FLOOR(RANDOM() * max_id) + 1;
        IF user1 <> user2 THEN
            BEGIN
                INSERT INTO friends (user_id, friend_id)
                VALUES (LEAST(user1, user2), GREATEST(user1, user2))
                ON CONFLICT DO NOTHING;
                counter := counter + 1;
            EXCEPTION WHEN uniqueViolation THEN
                END;
        END IF;
    END LOOP;
END $$;

```

❓ How would you query these tables to find the friends of a user and friends of friends and friends of friends etc ...

```
WITH RECURSIVE friends_cte(id, jumps) AS (
    SELECT friend_id, 1
    FROM friends
    WHERE user_id = 1
    UNION ALL
    SELECT f.friend_id, fc.jumps + 1
    FROM friends f
    INNER JOIN friends_cte fc ON f.user_id = fc.id
    WHERE fc.jumps < 3
)
SELECT id, jumps
FROM friends_cte;
```

This is pretty unintuitive 😅 let's try a different solution!

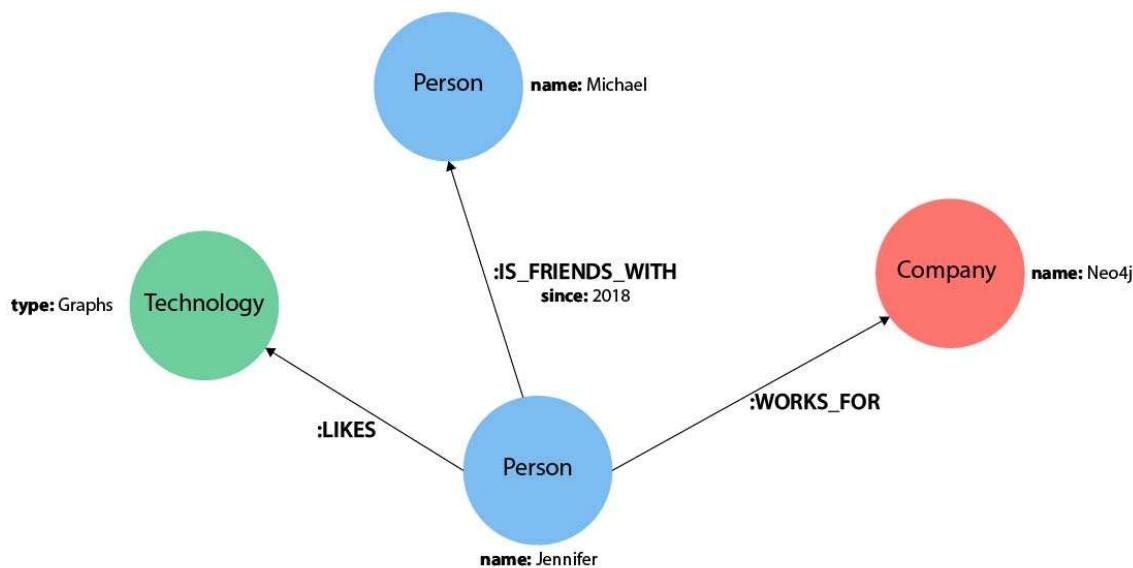
## 2 Graph Databases

e.g. Neo4j, Dgraph

**The easiest way to know when to use graph databases is when the relationships are as important as the data itself!**

## Graph model components

- **Vertices** (also known as nodes or entities)
- **Edges** (also known as relationships or arcs)
- **Properties** (extra information attached to the main two components)

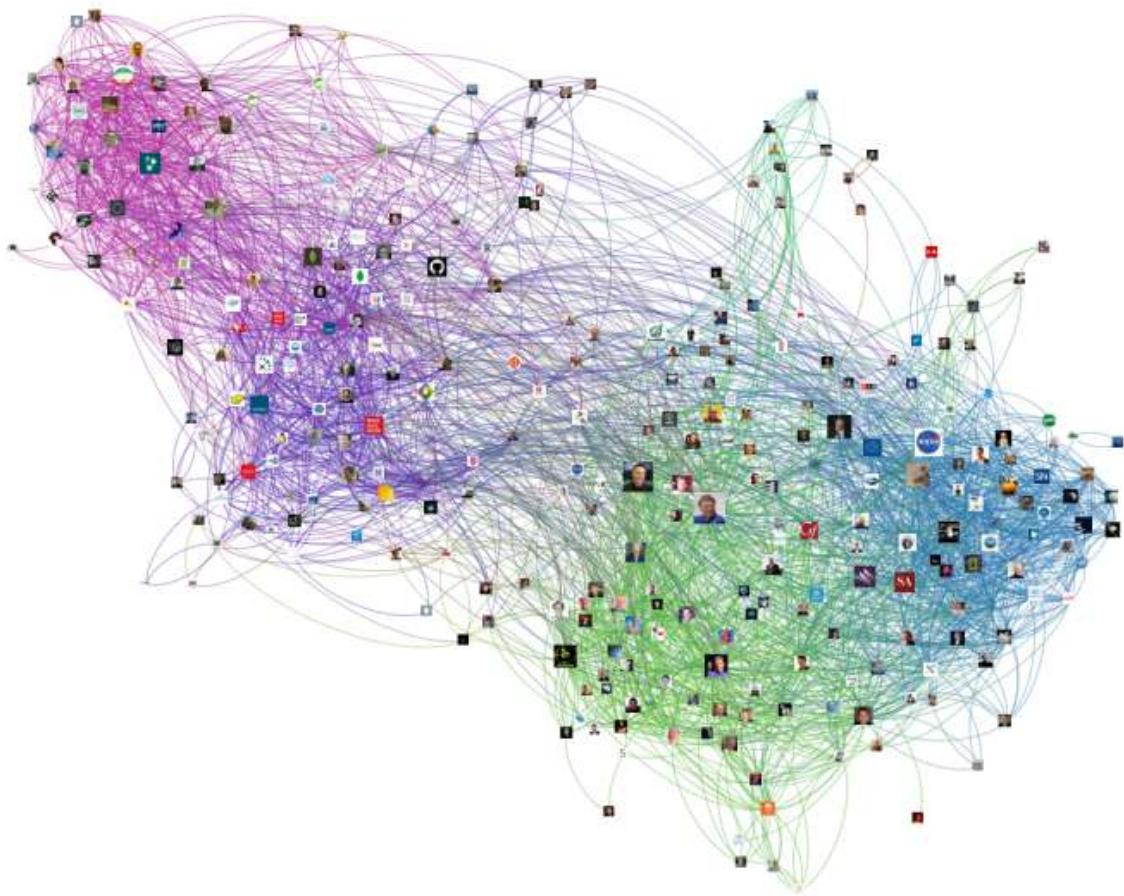


❓ For the following data what component do you think is appropriate?

- Users
- Date of Birth
- Which other users they follow
- Posts
- Likes
- Which users like the posts

## A few main usecases

### Social media and Social network graphs



## Fraud Detection

Graph representations make it very clear **a pattern** is out of place

Lots of fraudulent transactions involve **many relationships** for example tracing money via many bank accounts!

More recently **graph machine learning** has improved fraud detection even more!

## 360 Customer view

A newer use case which is gaining a lot of popularity

The idea is **centralizing a customers data** along with bringing the important information as relationships can lead to much richer understanding of a customer and their needs!

## Supply chain

Tracing goods through a supply chain is the perfect time to use a graph as it is a series of important relationships!

Become very popular with the rise of responsibility for your entire supply chain **legally and morally**

## 👉 Downsides of graph databases

- Scales poorly (jumping across servers really inefficient)
- Expensive
- Write operations slow
- Different language (easy enough to learn but restricts lots of people from your company)!

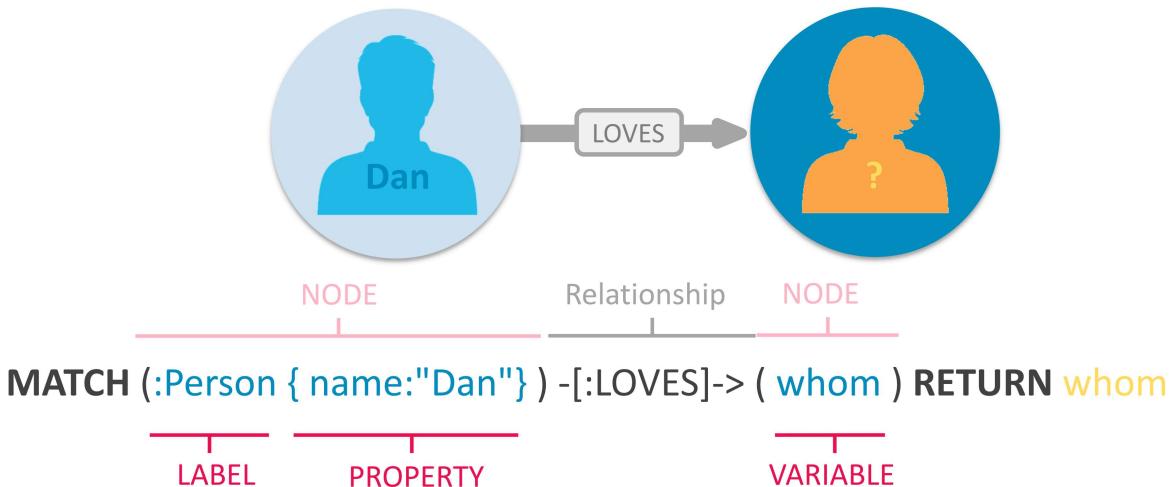
## 3 Neo4j

Neo4j is the **most popular graph database**

It uses its own query language **cypher**

It is **graph native** so efficient operations on graph problems!

## Cypher query language



## Testing neo4j with Docker

```
docker run \
  --name testneo4j \
  -p7474:7474 -p7687:7687 \
  -d \
  -v $HOME/neo4j/data:/data \
  -v $HOME/neo4j/logs:/logs \
  -v $HOME/neo4j/import:/var/lib/neo4j/import \
  -v $HOME/neo4j/plugins:/plugins \
  --env NEO4J_AUTH=neo4j/password \
  neo4j:latest
```

## Lets redo the task of modelling friends using neo4j

## **Creating the users and friends**

```

UNWIND range(0, 9999) AS counter
CREATE (:User {id: counter, name: 'User ' + counter, email: 'user' + counter +
'@example.com'});
UNWIND range(0, 19999) AS counter
WITH toInteger(rand() * 10000) AS user1, toInteger(rand() * 10000) AS user2
WHERE user1 <> user2
MERGE (u1:User {id: user1})
MERGE (u2:User {id: user2})
MERGE (u1)-[:FRIENDS]->(u2);

```

Same query as before to find friends **within two jumps of a user**

```

MATCH (user:User {id: 1})
MATCH (user)-[:FRIENDS*1..2]-(friend)
RETURN DISTINCT user, friend;

```

 **Shortest path** between two different users

```

MATCH (start:User {id: 1})
MATCH (end:User {id: 2})
MATCH path = shortestPath((start)-[:FRIENDS*]-(end))
RETURN path

```

## Info from python

```

from neo4j import GraphDatabase

class UserShortestPath:

    def __init__(self, uri, user, password):
        self.driver = GraphDatabase.driver(uri, auth=(user, password))

    def close(self):
        self.driver.close()

    def find_shortest_path(self, id1, id2):
        with self.driver.session() as session:
            path = session.execute_write(self._create_and_return_greeting, id1,
id2)
            path = path[0]["path"]
            for item in path:
                if item == "FRIENDS":
                    print(" "*5 + "|\\n" + " "*2 + "FRIENDS\\n" + " "*5 + "|")
                else:
                    print(item)

    @staticmethod
    def _create_and_return_greeting(tx, id1, id2):
        result = tx.run("MATCH (start:User {id: $id1})"
                      "MATCH (end:User {id: $id2})"
                      "MATCH path = shortestPath((start)-[:FRIENDS*]-(end))"
                      "RETURN path", id1=id1, id2=id2)
        return result.data()

if __name__ == "__main__":
    greeter = UserShortestPath("bolt://0.0.0.0:7687", "neo4j", "password")
    greeter.find_shortest_path(1, 2)
    greeter.close()

```

## Appendix

- Great lecture (<https://www.youtube.com/watch?v=yOYodfN84N4>)
- Case studies (<https://neo4j.com/use-cases/>)
- JanusGraph optimized for huge databases (<https://janusgraph.org/>)